



Writing MATLAB[®] C/MEX Code

Pascal Getreuer

What is MEX?

MEX-functions are programs written in C, C++, or Fortran code that are callable from MATLAB. We will focus on C.

MEX provides C functions to manipulate MATLAB variables, for example

C/MEX	Meaning
<code>mxCreateDoubleMatrix</code>	Create a 2D double array
<code>mxGetPr</code>	Get pointer to array data
<code>mxGetDimensions</code>	Get array dimensions
<code>mxGetClassID</code>	Determine variable's datatype

Words of Warning



My advice:

- 1 Try to optimize your M-code first
- 2 Consider porting the entire application to C/C++
- 3 Use MEX to substitute only the bottleneck M-functions

Getting Started

Hello, world!

hello.c

```
#include "mex.h" /* Always include this */

void mexFunction(int nlhs, mxArray *plhs[], /* Output variables */
                 int nrhs, const mxArray *prhs[]) /* Input variables */
{
    mexPrintf("Hello, world!\n"); /* Do something interesting */
    return;
}
```

Getting Started

On the MATLAB console, compile `hello.c` with

```
>> mex hello.c
```

Compiling requires that you have a C compiler and that MATLAB is configured to use it. Use “`mex -setup`” to change compiler settings.

Once the MEX-function is compiled, we can call it just like any M-file function:

```
>> hello  
Hello, world!
```

Beware that compiled MEX files might not be compatible between different platforms or different versions of MATLAB.

Inputs and Outputs

Let's take a closer look at the line

```
void mexFunction(int nlhs, mxArray *plhs[],  
                 int nrhs, const mxArray *prhs[])
```

“mxArray” is a type for representing a MATLAB variable, and the arguments are:

C/MEX	Meaning
nlhs	Number of output variables
plhs	Array of mxArray pointers to the output variables
nrhs	Number of input variables
prhs	Array of mxArray pointers to the input variables

Notation: “lhs” = “left-hand side” (output variables)
 “rhs” = “right-hand side” (input variables)

Inputs and Outputs

Let's take a closer look at the line

```
void mexFunction(int nlhs, mxArray *plhs[],  
                 int nrhs, const mxArray *prhs[])
```

“mxArray” is a type for representing a MATLAB variable, and the arguments are:

C/MEX	M-code equivalent
nlhs	nargout
plhs	varargout
nrhs	nargin
prhs	varargin

Notation: “lhs” = “left-hand side” (output variables)
 “rhs” = “right-hand side” (input variables)

Inputs and Outputs

Suppose our MEX-function is called as

```
[X, Y] = mymexfun(A, B, C)
```

Then `mexFunction`

```
void mexFunction(int nlhs, mxArray *plhs[],  
                 int nrhs, const mxArray *prhs[])
```

receives the following information:

`nlhs = 2`

`plhs[0]` points to X

`plhs[1]` points to Y

`nrhs = 3`

`prhs[0]` points to A

`prhs[1]` points to B

`prhs[2]` points to C

Inputs and Outputs

The output variables are initially unassigned; it is the responsibility of the MEX-function to create them. So for the example

```
[X, Y] = mymexfun(A, B, C)
```

it is our responsibility to create X and Y.

If `nlhs = 0`, the MEX-function is still allowed return one output variable, in which case `plhs[0]` represents the “ans” variable.

normalizecols

Objective: Given matrix A , construct matrix B according to

$$B_{m,n} = A_{m,n} / \|A_n\|_p$$

where $\|A_n\|_p$ is the ℓ^p norm of the n th column.

normalizecols.m

M-function implementation

```
function B = normalizecols(A, p)

if nargin < 2    % Was p not specified?
    p = 2;      % Set default value for p
end

[M, N] = size(A);
B = zeros(M, N);

for n = 1:N    % Build matrix B column-by-column
    B(:,n) = A(:,n) / norm(A(:,n), p);
end
```

normalizecols: Converting to MEX

We will convert `normalizecols.m` to MEX. Let's start with `hello.c` as a template...

`normalizecols.c`

MEX-function implementation

```
#include "mex.h" /* Always include this */

void mexFunction(int nlhs, mxArray *plhs[],          /* Outputs */
                 int nrhs, const mxArray *prhs[]) /* Inputs */
{
    /* TODO ... */
}
```

normalizecols: Converting to MEX

The first step is to figure out the calling syntax. Let's look at

```
B = normalizecols(A, p)
```

This calling syntax in MEX representation is

```
nlhs = 1
```

```
nrhs = 2
```

```
plhs[0] points to B
```

```
prhs[0] points to A
```

```
prhs[1] points to p
```

For clarity, it is a good idea to define

```
#define A_IN      prhs[0]
```

```
#define P_IN      prhs[1]
```

```
#define B_OUT     plhs[0]
```

normalizecols: Converting to MEX

What if `normalizedcols` is called without `p`?

```
B = normalizecols(A)
```

This calling syntax in MEX representation is

<code>nlhs = 1</code>	<code>nrhs = 1</code>
<code>plhs[0]</code> points to B	<code>prhs[0]</code> points to A
	<code>prhs[1]</code> doesn't exist

For clarity, it is a good idea to define

```
#define A_IN    prhs[0]  
#define P_IN    prhs[1]  
#define B_OUT    plhs[0]
```

normalizecols: Input Checking

Be on the defensive and check inputs. We can use `mexErrMsgTxt` to abort with an error message.

normalizecols.c

MEX-function implementation

```
#include "mex.h" /* Always include this */

#define A_IN      prhs[0]
#define P_IN      prhs[1]
#define B_OUT     plhs[0]

void mexFunction(int nlhs, mxArray *plhs[],          /* Outputs */
                  int nrhs, const mxArray *prhs[]) /* Inputs */
{
    if(nrhs < 1 || nrhs > 2)
        mexErrMsgTxt("Must have either 1 or 2 input arguments.");
    if(nlhs > 1)
        mexErrMsgTxt("Too many output arguments.");

    /* ... */
}
```

normalizecols: Input Checking

We should also verify the datatype of the input variables.

C/MEX	Meaning
<code>mxIsDouble(A_IN)</code>	True for a double array
<code>mxIsComplex(A_IN)</code>	True if array is complex
<code>mxIsSparse(A_IN)</code>	True if array is sparse
<code>mxGetNumberOfDimensions(A_IN)</code>	Number of array dimensions
<code>mxGetNumberOfElements(A_IN)</code>	Number of array elements

For simplicity, let's require that A is a real 2D full double matrix.

```
if(mxIsComplex(A_IN) || mxGetNumberOfDimensions(A_IN) != 2  
    || mxIsSparse(A_IN) || !mxIsDouble(A_IN))  
    mexErrMsgTxt("Sorry! A must be a real 2D full double matrix.");
```

normalizecols: Input Checking

We want to allow two calling syntaxes

```
B = normalizecols(A)      % nrhs == 1  (use p = 2.0)
B = normalizecols(A, p)   % nrhs == 2
```

We can determine which way `normalizecols` was called by checking `nrhs`. If `nrhs = 2`, then we should also check the datatype of `p`.

```
double p;

if(nrhs == 1)      /* Was p not specified? */
    p = 2.0;      /* Set default value for p */
else
    if(mxIsComplex(P_IN) || !mxIsDouble(P_IN)
       || mxGetNumberOfElements(P_IN) != 1)
        mexErrMsgTxt("p must be a double scalar.");
    else
        p = mxGetScalar(P_IN); /* Get the value of p */
```


normalizecols: Input Checking

normalizecols.c

MEX-function implementation

```
#include "mex.h" /* Always include this */

#define A_IN    prhs[0]
#define P_IN    prhs[1]
#define B_OUT   plhs[0]

void mexFunction(int nlhs, mxArray *plhs[],          /* Outputs */
                 int nrhs, const mxArray *prhs[]) /* Inputs */
{
    double p;

    /*** Check inputs ***/
    if(nrhs < 1 || nrhs > 2)
        mexErrMsgTxt("Must have either 1 or 2 input arguments.");
    if(nlhs > 1)
        mexErrMsgTxt("Too many output arguments.");

    if(mxIsComplex(A_IN) || mxGetNumberOfDimensions(A_IN) != 2
       || mxIsSparse(A_IN) || !mxIsDouble(A_IN))
        mexErrMsgTxt("Sorry! A must be a real 2D full double matrix.");

    if(nrhs == 1) /* Was p not specified? */
        p = 2.0; /* Set default value for p */
    else
        if(mxIsComplex(P_IN) || !mxIsDouble(P_IN)
           || mxGetNumberOfElements(P_IN) != 1)
            mexErrMsgTxt("p must be a double scalar.");
        else
            p = mxGetScalar(P_IN); /* Get the value of p */
}
```

normalizecols: Reading Input A

Now that the inputs are verified, we can safely interpret A as a real 2D full double matrix.

```
int M, N;  
double *A;  
  
/** Read matrix A */  
M = mxGetM(A_IN);      /* Get the dimensions of A */  
N = mxGetN(A_IN);  
A = mxGetPr(A_IN);     /* Get pointer to A's data */
```

The elements of A are stored contiguously in memory in column-major format,

$A[m + M*n]$ corresponds to $A(m+1,n+1)$

normalizecols: Creating Output B

Output variables must be created by the MEX-function. We can create B by

```
double *B;  
  
/** Create the output matrix **/  
B_OUT = mxCreateDoubleMatrix(M, N, mxREAL);  
B = mxGetPr(B_OUT); /* Get pointer to B's data */
```

The interface is now set up!

normalizecols: Interface Complete

normalizecols.c

MEX-function implementation

```
#include "mex.h" /* Always include this */

#define A_IN    prhs[0]
#define P_IN    prhs[1]
#define B_OUT   plhs[0]

void mexFunction(int nlhs, mxArray *plhs[],          /* Outputs */
                 int nrhs, const mxArray *prhs[]) /* Inputs */
{
    double p, *A, *B;
    int M, N;

    /** Check inputs **/
    /* ... */

    /** Read matrix A **/
    M = mxGetM(A_IN); /* Get the dimensions of A */
    N = mxGetN(A_IN);
    A = mxGetPr(A_IN); /* Get pointer to A's data */

    /** Create the output matrix **/
    B_OUT = mxCreateDoubleMatrix(M, N, mxREAL);
    B = mxGetPr(B_OUT); /* Get pointer to B's data */

    /** TODO: Do the computation itself */
    DoComputation(B, A, M, N, p);
}
```

normalizecols: DoComputation

All that remains is to code the computation itself.

```
#include <math.h>

void DoComputation(double *B, double *A, int M, int N, double p)
{
    double colnorm;
    int m, n;

    for(n = 0; n < N; n++)
    {
        /* Compute the norm of the nth column */
        for(m = 0, colnorm = 0.0; m < M; m++)
            colnorm += pow(A[m + M*n], p);

        colnorm = pow(fabs(colnorm), 1.0/p);

        /* Fill the nth column of B */
        for(m = 0; m < M; m++)
            B[m + M*n] = A[m + M*n]/colnorm;
    }
}
```

normalizecols: M-code vs. MEX

The MEX implementation is certainly more complicated than the M-function. So did our hard work pay off?

Runtime (s) with $p = 2.7$:

Input A	M-function	MEX-function	MEX-function*
1000×1000	0.2388	0.0963	0.0927
2000×2000	0.9479	0.3656	0.3599
3000×3000	2.1976	0.8896	0.8550
4000×4000	3.9033	1.5816	1.5128

*With minor optimizations

More Information

Please see [Writing MATLAB C/MEX Code](http://www.math.ucla.edu/~getreuer) on my webpage
www.math.ucla.edu/~getreuer



Tutorials

- [Writing MATLAB C/MEX Code](#)
MEX: Combine the power of MATLAB and C.
- [Writing Fast MATLAB Code](#)
Profiling, JIT, vectorization, and more.
- [Image Processing with MATLAB](#)
Reading and writing image files, basic operations, and filtering
- [TikZ for High-Quality L^AT_EX Pictures](#)